

B561 Advanced Database Concepts

Assignment 2

Fall 2021

Solutions

Dirk Van Gucht

This assignment relies on the lectures

- SQL Part 1 and SQL Part 2 (Pure SQL);
- Views;
- Relational Algebra (RA); and
- Joins and semijoins.

Furthermore, the lecture on the correspondence between Tuple Relational Calculus and SQL should help you to solve problems in this assignment.

To turn in your assignment, you will need to upload to Canvas a single file with name `assignment2.sql` which contains the necessary SQL statements that solve the problems in this assignment. The `assignment2.sql` file must be so that the AI's can run it in their PostgreSQL environment. You should use the `Assignment2Script.sql` file to construct the `assignment2.sql` file. (Note that the data to be used for this assignment is included in this file.) In addition, you will need to upload a separate `assignment2.txt` file that contains the results of running your queries. Finally, you need to upload a file `assignment2.pdf` that contains the solutions to the problems that require it.

The problems that need to be included in the `assignment2.sql` are marked with a blue bullet •. The problems that need to be included in the `assignment2.pdf` are marked with a red bullet •. (You should aim to use Latex to construct your .pdf file.)

Database schema and instances

For the problems in this assignment we will use the following database schema:¹

```
Person(pid, pname, city)
Company(cname, headquarter)
Skill(skill)
worksFor(pid, cname, salary)
companyLocation(cname, city)
personSkill(pid, skill)
hasManager(eid, mid)
Knows(pid1, pid2)
```

In this database we maintain a set of persons (**Person**), a set of companies (**Company**), and a set of (job) skills (**Skill**). The **pname** attribute in **Person** is the name of the person. The **city** attribute in **Person** specifies the city in which the person lives. The **cname** attribute in **Company** is the name of the company. The **headquarter** attribute in **Company** is the name of the city wherein the company has its headquarter. The **skill** attribute in **Skill** is the name of a (job) skill.

A person can work for at most one company. This information is maintained in the **worksFor** relation. (We permit that a person does not work for any company.) The **salary** attribute in **worksFor** specifies the salary made by the person.

The **city** attribute in **companyLocation** indicates a city in which the company is located. (Companies may be located in multiple cities.)

A person can have multiple job skills. This information is maintained in the **personSkill** relation. A job skill can be the job skill of multiple persons. (A person may not have any job skills, and a job skill may have no persons with that skill.)

A pair (e, m) in **hasManager** indicates that person e has person m as one of his or her managers. We permit that an employee has multiple managers and that a manager may manage multiple employees. (It is possible that an employee has no manager and that an employee is

¹The primary key, which may consist of one or more attributes, of each of these relations is underlined.

not a manager.) We further require that an employee and his or her managers must work for the same company.

The relation **Knows** maintains a set of pairs (p_1, p_2) where p_1 and p_2 are pids of persons. The pair (p_1, p_2) indicates that the person with pid p_1 knows the person with pid p_2 . We do not assume that the relation **Knows** is symmetric: it is possible that (p_1, p_2) is in the relation but that (p_2, p_1) is not.

The domain for the attributes **pid**, **pid1**, **pid2**, **salary**, **eid**, and **mid** is **integer**. The domain for all other attributes is **text**.

We assume the following foreign key constraints:

- **pid** is a foreign key in **worksFor** referencing the primary key **pid** in **Person**;
- **cname** is a foreign key in **worksFor** referencing the primary key **cname** in **Company**;
- **cname** is a foreign key in **companyLocation** referencing the primary key **cname** in **Company**;
- **pid** is a foreign key in **personSkill** referencing the primary key **pid** in **Person**;
- **skill** is a foreign key in **personSkill** referencing the primary key **skill** in **Skill**;
- **eid** is a foreign key in **hasManager** referencing the primary key **pid** in **Person**;
- **mid** is a foreign key in **hasManager** referencing the primary key **pid** in **Person**;
- **pid1** is a foreign key in **Knows** referencing the primary key **pid** in **Person**; and
- **pid2** is a foreign key in **Knows** referencing the primary key **pid** in **Person**

The file **Assignment2Script.sql** contains the data supplied for this assignment.

Pure SQL and RA SQL

In this assignment, we distinguish between Pure SQL and RA SQL. Below we list the **only** features that are allowed in Pure SQL and in RA SQL.

In particular notice that

- JOIN, NATURAL JOIN, and CROSS JOIN are **not** allowed in Pure SQL.
- The predicates [NOT] IN, SOME, ALL, [NOT] EXISTS are **not** allowed in RA SQL.

The only features allowed in Pure SQL

SELECT ... FROM ... WHERE
WITH ...
UNION, INTERSECT, EXCEPT operations
EXISTS and NOT EXISTS predicates
IN and NOT IN predicates
ALL and SOME predicates
VIEWS that can only use the above RA SQL features

The only features allowed in Pure SQL features

SELECT ... FROM ... WHERE
WITH ...
UNION, INTERSECT, EXCEPT operations
JOIN ... ON ..., NATURAL JOIN, and CROSS JOIN operations
VIEWS that can only use the above RA SQL features
commas in the FROM clause are **not** allowed

1 Formulating queries in Pure SQL with and without set predicates

An important consideration in formulating queries is to contemplate how they can be written in different, but equivalent, ways. In fact, this is an aspect of programming in general and, as can be expected, is also true for SQL. A learning outcome of this course is to acquire skills for writing queries in different ways. One of the main motivations for this is to learn that different formulations of the same query can dramatically impact performance, sometimes by orders of magnitude.

For the problems in this section, you will need to formulate queries in Pure SQL with and without set predicates. You can use the SQL operators `INTERSECT`, `UNION`, and `EXCEPT`, unless otherwise specified. You are however allowed and encouraged to use views including temporary and user-defined views.

To illustrate what you need to do, consider the following example.

Example 1 *Consider the query “Find the pid and name of each person who (a) works for a company located in Bloomington and (b) knows a person who lives in Chicago.”*

- (a) *Formulate this query in Pure SQL by only using the `EXISTS` or `NOT EXISTS` set predicates. You can not use the set operations `INTERSECT`, `UNION`, and `EXCEPT`.*

A possible solution is

```
select p.pid, p.pname
from   Person p
where  exists (select 1
               from   worksFor w
               where  p.pid = w.pid and
                     exists (select 1
                             from   companyLocation c
                             where  w.cname = c.cname and c.city = 'Bloomington')) and
        exists (select
               from   Knows k
               where  p.pid = k.pid1 and
                     exists (select 1
                             from   Person p2
                             where  k.pid2 = p2.pid and
                                     p2.city = 'Chicago'));
```

- (b) *Formulate this query in Pure SQL by only using the IN, NOT IN, SOME, or ALL set membership predicates. You can not use the set operations INTERSECT, UNION, and EXCEPT.*

A possible solution is

```
select p.pid, p.pname
from   Person p
where  p.pid in (select w.pid
                  from   worksFor w
                  where  w.cname in (select c.cname
                                     from   companyLocation c
                                     where  c.city = 'Bloomington')) and
        p.pid in (select k.pid1
                  from   Knows k
                  where  k.pid2 in (select p2.pid
                                     from   Person p2
                                     where  p2.city = 'Chicago'));
```

Another possible solution using the SOME and IN predicates

```
select p.pid, p.pname
from   Person p
where  p.pid = some (select w.pid
                     from   worksFor w
                     where  w.cname = some (select c.cname
                                             from   companyLocation c
                                             where  c.city = 'Bloomington')) and
        p.pid in (select k.pid1
                  from   Knows k
                  where  k.pid2 in (select p2.pid
                                     from   Person p2
                                     where  p2.city = 'Chicago'));
```

- (c) *Formulate this query in Pure SQL without using set predicates. A possible solution is*

```
select p.pid, p.pname
from   Person p, worksFor w, companyLocation c
where  p.pid = w.pid and
        w.cname = c.cname and
        c.city = 'Bloomington'
intersect
select p1.pid, p1.pname
from   Person p1, Knows k, Person p2
where  k.pid1 = p1.pid and
        k.pid2 = p2.pid and
        p2.city = 'Chicago';
```

We now turn to the problems for this section.

1. Consider the query “*Find the pid and name of each person who works for Google and who has a strictly higher salary than some other person who he or she knows and who also works for Google.*”
 - (a) • Formulate this query in Pure SQL by only using the EXISTS or NOT EXISTS set predicates.
 - (b) • Formulate this query in Pure SQL by only using the IN, NOT IN, SOME, or ALL set membership predicates.
 - (c) • Formulate this query in Pure SQL without using set predicates.
2. Consider the query “*Find the cname of each company with head-quarter in Cupertino, but that is not located in Indianapolis, along with the pid, name, and salary of each person who works for that company and who has the next-to-lowest salary at that company.*”²
 - (a) • Formulate this query in Pure SQL by only using the EXISTS or NOT EXISTS set predicates. You can not use the set operations INTERSECT, UNION, and EXCEPT.
 - (b) • Formulate this query in Pure SQL by only using the IN, NOT IN, SOME, or ALL set membership predicates. You can not use the set operations INTERSECT, UNION, and EXCEPT.
 - (c) • Formulate this query in Pure SQL without using set predicates.
3. Consider the query “*Find each (c,p) pair where c is the cname of a company and p is the pid of a person who works for that company and who is known by all other persons who work for that company.*”
 - (a) • Formulate this query in Pure SQL by only using the EXISTS or NOT EXISTS set predicates. You can not use the set operations INTERSECT, UNION, and EXCEPT.

²By definition, a salary s is next-to-lowest if there exists salary s_1 with $s_1 < s$, but there do not exist two salaries s_1 and s_2 with $s_2 < s_1 < s$.

- (b) • Formulate this query in Pure SQL by only using the `IN`, `NOT IN`, `SOME`, or `ALL` set membership predicates. You can not use the set operations `INTERSECT`, `UNION`, and `EXCEPT`.
 - (c) • Formulate this query in Pure SQL without using set predicates.
4. Consider the query “*Find each skill that is not a jobskill of any person who works for Yahoo or for Netflix.*”
- (a) • Formulate this query in Pure SQL using subqueries and set predicates. You can not use the set operations `INTERSECT`, `UNION`, and `EXCEPT`.
 - (b) • Formulate this query in Pure SQL without using predicates.
5. Consider the query “*Find the pid and name of each person who manages all-but-one person who works for Google.*”
- (a) • Formulate this query in Pure SQL using subqueries and set predicates. You can not use the set operations `INTERSECT`, `UNION`, and `EXCEPT`.
 - (b) • Formulate this query in Pure SQL without using set predicates.

2 Formulating queries in Relational Algebra and RA SQL

Reconsider the queries in Section 1. The goal of the problems in this section is to formulate these queries in Relational Algebra in standard notation and in RA SQL.

There are some further restrictions:

- You can only use **WHERE** clauses that use conditions involving constants. For example conditions of the form “ $t.A \theta 'a$ ” are allowed, but conditions of the form “ $'t.A \theta s.B$ ” are not allowed. The latter conditions can be absorbed in JOIN operations in the **FROM** clause.
- You can not use commas in any **FROM** clause. Rather, you should use JOIN operations.

You can use the following letters, or indexed letters, to denote relation names in RA expressions:

P, P_1, P_2, \dots	Person
C, C_1, C_2, \dots	Company
S, S_1, S_2, \dots	Skill
W, W_1, W_2, \dots	worksFor
cL, cL_1, cL_2, \dots	companyLocation
pS, pS_1, pS_2, \dots	personSkill
M, M_1, M_2, \dots	hasManager
K, K_1, K_2, \dots	Knows

To illustrate what you need to do reconsider the query in Example 1 in Section 1.

Example 2 Consider the query “Find the pid and name of each person who (a) works for a company located in Bloomington and (b) knows a person who lives in Chicago.”

(a) Formulate this query in Relational Algebra in standard notation.

A possible solution is

$$\pi_{pid, pname}(Person \bowtie worksFor \bowtie \pi_{cname}(\sigma_{city=Bloomington}(companyLocation))) \cap \pi_{Person_1.pid, Person_1.pname}(Person_1 \bowtie_{Person_1.pid=pid1} Knows \bowtie_{pid2=Person_2.pid} \pi_{Person_2.pid}(\sigma_{city=Chicago}(Person_2)))$$

If we use the letters in the above box, this expression becomes more succinct:

$$\pi_{pid, pname}(P \bowtie W \bowtie \pi_{cname}(\sigma_{city=\mathbf{Bloomington}}(cL))) \cap \pi_{P_1.pid, P_1.pname}(P_1 \bowtie_{P_1.pid=pid1} K \bowtie_{pid2=P_2.pid} \pi_{P_2.pid}(\sigma_{city=\mathbf{Chicago}}(P_2)))$$

You are also allowed to introduce letters that denote expressions. For example, let E and F denote the expression

$$\pi_{pid, pname}(P \bowtie W \bowtie \pi_{cname}(\sigma_{city=\mathbf{Bloomington}}(cL)))$$

and

$$\pi_{P_1.pid, P_1.pname}(P_1 \bowtie_{P_1.pid=pid1} K \bowtie_{pid2=P_2.pid} \pi_{P_2.pid}(\sigma_{city=\mathbf{Chicago}}(P_2))),$$

respectively. Then we can write the solution as follows:

$$E \cap F.$$

(b) Formulate this query in RA SQL.

A possible solution is

```
select pid, pname
from   Person
      NATURAL JOIN worksFor
      NATURAL JOIN (select cname
                    from   companyLocation
                    where  city = 'Bloomington') C
INTERSECT
select P1.pid, P1.pname
from   Person P1
      JOIN Knows ON (P1.pid = pid1)
      JOIN (select pid
            from   Person
            where  city = 'Chicago') P2 ON (pid2 = P2.pid)
order by 1,2;
```

Observe that the **WHERE** clauses only use conditions involving constants.

We now turn to the problems in this section.

6. Reconsider Problem 1. Find the pid and name of each person who works for Google and who has a strictly higher salary than some other person who he or she knows and who also works for Google.

- (a) • Formulate this query in Relational Algebra in standard notation.

Let W^G denote the expression

$$\pi_{pid, salary}(\sigma_{cname=\mathbf{Google}}(W)).$$

Then a possible RA expression for the query is

$$\pi_{P.pid, P.pname}(P \bowtie_{P.pid=W_1^G.pid} W_1^G \bowtie_{W_1^G.salary > W_2^G.salary} W_2^G \bowtie_{W_1^G.pid=K.pid1 \wedge W_2^G.pid=K.pid2} K)$$

- (b) • Formulate this query in RA SQL.

```
with   worksForGoogle as (select pid, salary from worksfor where cname = 'Google')
select distinct p.pid, p.pname
from   Person p
       JOIN worksforGoogle w1 ON (p.pid = w1.pid)
       JOIN worksforGoogle w2 ON (w1.salary > w2.salary)
       JOIN Knows ON (w1.pid = pid1 and w2.pid = pid2);
```

7. Reconsider Problem 2. Find the cname of each company with headquarter in Bloomington, but not located in Indianapolis, along with the pid, name, and salary of each person who works for that company and who has the next-to-lowest salary at that company.

- (a) • Formulate this query in Relational Algebra in standard notation.

Consider the following RA expressions:

$$\begin{aligned} C &= \pi_{cname}(\sigma_{headquarter=\mathbf{Cupertino}}(Company)) - \pi_{cname}(\sigma_{city=\mathbf{Indianapolis}}(companyLocation)) \\ E_1 &= P \bowtie_{P_1.cname=P.cname \wedge P_1.salary < P.salary} P_1 \\ E_2 &= E_1 \bowtie_{P_2.cname=P.cname \wedge P_2.salary < P_1.salary \text{ and } P_1.salary < P.salary} P_2 \end{aligned}$$

Then a possible RA expression for the query is

$$E_1 - \pi_{P.cname, P.pid, P.pname, P.salary}(E_2).$$

- (b) • Formulate this query in RA SQL.

```

with company as (select  cname
                   from    company
                   where    headquarter = 'Cupertino'
                   except
                   select  cname
                   from    companyLocation
                   where    city = 'Indianapolis'),
person as (select c.cname, p.pid, p.pname, w.salary
              from  person p
                  JOIN worksfor w ON (p.pid = w.pid)
                  JOIN company c ON (w.cname = c.cname))

select p.*
from   person p
      JOIN person p1 ON p1.cname = p.cname and p1.salary < p.salary
except
select p.*
from   person p
      JOIN person p1 ON (p1.cname = p.cname and p1.salary < p.salary)
      JOIN person p2 ON (p2.cname = p.cname and
                        p2.salary < p1.salary and
                        p1.salary < p.salary);

```

8. Reconsider Problem 3. Find each (c, p) pair where c is the cname of a company and p is the pid of a person who works for that company and who is known by all other persons who work for that company.

- (a) • Formulate this query in Relational Algebra in standard notation.

Let E be the expression

$$W \bowtie_{W.cname=W_1.cname \wedge W.pid \neq W_1.pid} W_1$$

Then a possible RA expression for the query is

$$\pi_{cname, pid}(W) - \pi_{W.cname, W.pid}(\pi_{W.cname, W.pid, W_1.cname, W_2.pid}(E) - \pi_{W.cname, W.pid, W_1.cname, W_2.pid}(E \bowtie_{W_1.pid=K.pid1 \wedge W.pid=K.pid2} K))$$

- (b) • Formulate this query in RA SQL.

```

select cname, pid
from   worksFor
except
(
select q.cname, q.pid
from
  (select w.cname, w.pid, w1.cname as w1cname, w1.pid as w1pid
   from   worksFor w
        JOIN worksFor w1 ON (w1.cname = w.cname and w1.pid <> w.pid)
   except
   select w.cname, w.pid, w1.cname as w1cname, w1.pid as w1pid
   from   worksFor w
        JOIN worksFor w1 ON (w1.cname = w.cname and w1.pid <> w.pid)
        JOIN Knows k ON (w1.pid = pid1 and w.pid = pid2)) q
)
order by 1,2;

```

9. Reconsider Problem 4. Find each skill that is not a jobskill of any person who works for Yahoo or for Netflix.

- (a) • Formulate this query in Relational Algebra in standard notation.

$$\pi_{skill}(S) - \pi_{S.skill}(S \bowtie_{S.skill=pS.skill} pS \bowtie_{pS.pid=W.pid \wedge (W.cname=Yahoo \vee W.cname=Netflix} W)$$

- (b) • Formulate this query in RA SQL.

```

select skill
from   Skill
except
select s.skill
from   Skill s
      JOIN personSkill ps ON (s.skill = ps.skill)
      JOIN worksFor w ON (ps.pid = w.pid and (w.cname = 'Yahoo' or w.cname = 'Netflix'));

```

10. Reconsider Problem 5. Find the pid and name of each person who manages all-but-one person who works for Google.

- (a) • Formulate this query in Relational Algebra in standard notation.

Consider the following expression

$$P^G = \pi_{P.pname, P.pid}(P \bowtie_{P.pid=W.pid} \sigma_{cname=Google}(W))$$

$$E_1 = \pi_{P^G.*, P_1^G.*, P_2^G.*}(P^G \bowtie_{P^G.pid=hm.pid} hM \bowtie_{P_1^G.pid=hm.eid} P_1^G \times P_2^G)$$

$$E_2 = \pi_{P^G.*, P_2^G.*, P_1^G.*}(E_1)$$

Then a possible RA expression for the query is

$$P^G - \pi_{P^G.pid, P^G.pname}(P^G \times (P_1^G \bowtie_{P_1^G.pid \neq P_2^G.pid} P_2^G) - (E_1 \cup E_2))$$

(b) • Formulate this query in RA SQL.

```

with personGoogle as (select p.pid, p.pname
                        from   person p JOIN worksFor w ON (p.pid = w.pid)
                        where  w.cname = 'Google')

select p.pid, pname
from   personGoogle p
except
(select q.ppuid, q.ppname
from
  (select p.pid as ppuid, p.pname as ppname, p1.*, p2.*
   from   personGoogle p
         CROSS JOIN
         (personGoogle p1 JOIN personGoogle p2 on p1.pid <> p2.pid)
  except
  (select p.*, p1.*, p2.*
   from   personGoogle p
         JOIN hasmanager hm ON (p.pid = hm.mid)
         JOIN personGoogle p1 ON (p1.pid = hm.eid)
         CROSS JOIN personGoogle p2
  union
  select p.*, p2.*, p1.*
   from   personGoogle p
         JOIN hasmanager hm ON (p.pid = hm.mid)
         JOIN personGoogle p1 ON (p1.pid = hm.eid)
         CROSS JOIN personGoogle p2)) q
);

```

3 Formulating constraints using Relational Algebra

Recall that it is possible to express constraints in TRC and as boolean SQL queries. It is also possible to write constraints using RA expressions. Let E , F , and G denote RA expressions. Then we can write RA expression comparisons that express constraints:

- $E \neq \emptyset$ which is true if E evaluates to a **non-empty** relation
- $E = \emptyset$ which is true if E evaluates to the **empty** relation
- $F \subseteq G$ which is true if F evaluates to a relation that is a **subset** of the relation obtained from G
- $F \not\subseteq G$ which is true if F evaluates to a relation that is **not** a **subset** of the relation obtained from G

Here are some examples of writing constraints in this manner.

Example 3 “Some person works for Google.” *This constraint can be written as follows:*

$$\pi_{pid}(\sigma_{cname=Google}(worksFor)) \neq \emptyset.$$

Indeed, the RA expression

$$\pi_{pid}(\sigma_{cname=Google}(worksFor))$$

computes the set of all person pids who work for Google. If this set is not empty then there are indeed persons who work for Google.

Incidentally, the constraint “No one works for Google” can be written as follows:

$$\pi_{pid}(\sigma_{cname=Google}(worksFor)) = \emptyset.$$

Example 4 Each person has at least two skills. *This constraint can be written as follows:*

$$\pi_{pid}(P) \subseteq \pi_{pS_1.pid}(\sigma_{pS_1.skill \neq pS_2.skill}(pS_1 \bowtie_{pS_1.pid=pS_2.pid} pS_2)).$$

Indeed,

$$\pi_{pid}(P)$$

computes the set of all person pids and

$$\pi_{pS_1.pid}(\sigma_{pS_1.skill \neq pS_2.skill}(pS_1 \bowtie_{pS_1.pid=pS_2.pid} pS_2))$$

computes the set of all pids of persons who have at least two skills. When the first set is contained in the second we must have that each person has at least two skills.

Incidentally, the constraint “Some person has fewer than two skills” can be written as follows:

$$\pi_{pid}(P) \not\subseteq \pi_{pS_1.pid}(\sigma_{pS_1.skill \neq pS_2.skill}(pS_1 \bowtie_{pS_1.pid=pS_2.pid} pS_2)).$$

We now turn to the problems in this section.

Formulate each of the following constraints using RA expressions as illustrated in Example 3 and Example 4.

11. • Some person has a salary that is strictly lower than that of each of the persons who he or she manages.

$$\pi_{eid}(M) \not\subseteq \pi_{M.eid}(\pi_{W.pid, W_1.pid}(W \bowtie_{W.salary \geq W_1.salary} W_1) \cap M).$$

12. • No person knows all persons who work for Google.

$$\pi_{pid}(P) \subseteq \pi_{P.pid}(\pi_{pid}(P) \times \pi_{pid}(\sigma_{cname=Google}(W)) - Knows).$$

13. • Each person knows all of his or her managers.

$$M \subseteq K.$$

14. • Each employee and his or her managers work for the same company.

$$\pi_{W_1.pid, W_2.pid}(W_1 \bowtie_{W_1.cname \neq W_2.cname} W_2) \cap M = \emptyset.$$

15. • The attribute pid is a primary key of the Person relation.

$$P_1 \bowtie_{P_1.pid=P_2.pid \wedge \neg(P_1.pname=P_2.pname \wedge P_1.city=P_2.city)} P_2 = \emptyset.$$

4 Formulating queries in SQL using views

Formulate the following views and queries in SQL. You are allowed to combine the features of both Pure SQL and RA SQL.

16. • Create a view **Triangle** that contains each triple of pids of different persons (p_1, p_2, p_3) that mutually know each other.
Then test your view.
17. • Define a parameterized view **SalaryBelow**(cname text, salary integer) that returns, for a given company identified by **cname** and a given **salary** value, the subrelation of **Person** of persons who work for company **cname** and whose salary is strictly below **salary**.
Test your view for the parameter values ('IBM',60000), ('IBM',50000), and ('Apple',65000).
18. • Define a parameterized view **KnowsPersonAtCompany**(p integer, c text) that returns for a person with pid p the subrelation of **Person** of persons who p knows and who work for the company with **cname** c .
Test your view for the parameters (1001, 'Amazon'), (1001,'Apple'), and (1015,'Netflix').
19. • Define a parameterized view **KnownByPersonAtCompany**(p integer, c text) that returns the subrelation of **Person** of persons who know the person with pid p and who work for the company with **cname** c .
Test your view for the parameters (1001, 'Amazon'), (1001,'Apple'), and (1015,'Netflix').
20. • Let $Tree(parent : integer, child : integer)$ be a rooted parent-child tree. So a pair (n, m) in $Tree$ indicates that node n is a parent of node m . The **sameGeneration**(n1, n2) binary relation is inductively defined using the following two rules:
 - If n is a node in T , then the pair (n, n) is in the **sameGeneration** relation. (**Base rule**)

- If n_1 is the parent of m_1 in *Tree* and n_2 is the parent of m_2 in *Tree* and (n_1, n_2) is a pair in the `sameGeneration` relation then (m_1, m_2) is a pair in the `sameGeneration` relation.
(Inductive Rule)

Write a [recursive view](#) for the `sameGeneration` relation.

Test your view.